

FINAL EXAM

ISyE6420

Fall 2020

Released November 24, 12:00pm – due December 8, 11:55pm. This exam is not proctored and not time limited except the due date. Late submissions will not be accepted.

Use of all available electronic and printed resources is allowed except direct communication that violates Georgia Tech Academic Integrity Rules.

Name Chen-Yang(Jim), Liu

Problem	1	2	3	Total
Score	/33	/33	/34	/100

ISyE6420_final

December 9, 2020

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[2]: import arviz as az
import pymc3 as pm
```

1 Vasoconstriction

The data give the presence or absence ($y_i = 1$ or 0) of vasoconstriction in the skin of the fingers following inhalation of a certain volume of air (v_i) at a certain average rate (r_i). Total number of records is 39. The candidate models for analyzing the relationship are the usual logit, probit, cloglog, loglog, and cauchyit models.

```
[3]: y_1 = np.array([1,1,1,1,1,1,0,0,0,0,0,0,0,1,1,1,1,1,
    0,1,0,0,0,0,1,0,1,0,1,0,1,0,0,1,1,1,0,0,1])
v_1 = np.array([3.7, 3.5, 1.25, 0.75, 0.8, 0.7, 0.6, 1.1, 0.9, 0.9, 0.8, 0.55, 0.
    →6, 1.4, 0.75, 2.3, 3.2, 0.85, 1.7,
    1.8, 0.4, 0.95, 1.35, 1.5, 1.6, 0.6, 1.8, 0.95, 1.9, 1.6, 2.7, 2.
    →35, 1.1, 1.1, 1.2, 0.8, 0.95, 0.75, 1.3])
r_1 = np.array([0.825, 1.09, 2.5, 1.5, 3.2, 3.5, 0.75, 1.7, 0.75, 0.45, 0.57, 2.
    →75, 3, 2.33, 3.75, 1.64, 1.6, 1.415,
    1.06, 1.8, 2, 1.36, 1.35, 1.36, 1.78, 1.5, 1.5, 1.9, 0.95, 0.4, 0.75, 0.3, 1.83,
    →2.2, 2, 3.33, 1.9, 1.9, 1.625])
```

1.1 (a) Transform covariates v and r as $x_1 = \log(10 \times v)$, $x_2 = \log(10 \times r)$.

```
[4]: x_1 = np.log(10 * v_1)
x_1
```

```
[4]: array([3.61091791, 3.55534806, 2.52572864, 2.01490302, 2.07944154,
    1.94591015, 1.79175947, 2.39789527, 2.19722458, 2.19722458,
    2.07944154, 1.70474809, 1.79175947, 2.63905733, 2.01490302,
```

```
3.13549422, 3.4657359 , 2.14006616, 2.83321334, 2.89037176,
1.38629436, 2.2512918 , 2.60268969, 2.7080502 , 2.77258872,
1.79175947, 2.89037176, 2.2512918 , 2.94443898, 2.77258872,
3.29583687, 3.15700042, 2.39789527, 2.39789527, 2.48490665,
2.07944154, 2.2512918 , 2.01490302, 2.56494936])
```

```
[5]: x_2 = np.log(10 * r_1)
x_2
```

```
[5]: array([2.1102132 , 2.38876279, 3.21887582, 2.7080502 , 3.4657359 ,
3.55534806, 2.01490302, 2.83321334, 2.01490302, 1.5040774 ,
1.74046617, 3.314186 , 3.40119738, 3.14845336, 3.62434093,
2.79728133, 2.77258872, 2.64971462, 2.360854 , 2.89037176,
2.99573227, 2.61006979, 2.60268969, 2.61006979, 2.87919846,
2.7080502 , 2.7080502 , 2.94443898, 2.2512918 , 1.38629436,
2.01490302, 1.09861229, 2.90690106, 3.09104245, 2.99573227,
3.5055574 , 2.94443898, 2.94443898, 2.78809291])
```

```
[6]: df = pd.DataFrame({"vasoconstriction": y_1, "air_log": x_1, "rate_log": x_2})
df.head()
```

```
[6]:
```

	vasoconstriction	air_log	rate_log
0	1	3.610918	2.110213
1	1	3.555348	2.388763
2	1	2.525729	3.218876
3	1	2.014903	2.708050
4	1	2.079442	3.465736

1.2 (b) Estimate posterior means for coefficients in the logit model. Use noninformative priors on all coefficients.

```
[7]: names = df.index.values
N = len(names)
dims={
    "air_log": ["developer"],
    "rate_log": ["developer"]
}
```

```
[8]: '''
pymc3 priors are
default_regressor_prior = Normal.dist(mu=0, tau=1.0E-6) and
default_intercept_prior = Flat.dist()
'''
with pm.Model() as logistic_model:
    pm.glm.linear.GLM(y=df["vasoconstriction"], x= df[["air_log", "rate_log"]],
    →intercept=True,
```

```
family=pm.glm.families.Binomial())
trace_log = pm.sample(5000, tune=5000, init='adapt_diag')
posterior_predictive = pm.sample_posterior_predictive(trace_log)
```

Auto-assigning NUTS sampler...
 Initializing NUTS using adapt_diag...
 Multiprocess sampling (4 chains in 4 jobs)
 NUTS: [rate_log, air_log, Intercept]

Sampling 4 chains for 5_000 tune and 5_000 draw iterations (20_000 + 20_000 draws total) took 54 seconds.
 There was 1 divergence after tuning. Increase `target_accept` or reparameterize.
 There were 26 divergences after tuning. Increase `target_accept` or reparameterize.
 The acceptance probability does not match the target. It is 0.6210488848956406, but should be close to 0.8. Try to increase the number of tuning steps.
 There were 2 divergences after tuning. Increase `target_accept` or reparameterize.
 The acceptance probability does not match the target. It is 0.7165668834870483, but should be close to 0.8. Try to increase the number of tuning steps.
 The number of effective samples is smaller than 10% for some parameters.

```
[9]: az.summary(trace_log)
```

```
[9]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_mean	\
Intercept	-30.883	10.584	-50.062	-12.367	0.408	0.289	672.0	
air_log	6.269	2.174	2.536	10.256	0.084	0.059	671.0	
rate_log	5.595	2.021	2.366	9.520	0.073	0.051	776.0	

	ess_sd	ess_bulk	ess_tail	r_hat
Intercept	672.0	519.0	322.0	1.01
air_log	671.0	523.0	336.0	1.01
rate_log	776.0	630.0	789.0	1.01

The means of intercept, air_log and rate_log are -30.883, 6.269 and 5.595, respectively.

1.3 (c) For a subject with $v = r = 1.5$, find the probability of vasoconstriction.

$$P(\text{vasoconstriction}=1) = p = \frac{1}{1 + e^{-(\text{intercept} + 6.461 \times \text{air}_{\log} + 5.719 \times \text{rate}_{\log})}}$$

```
[42]: prob_v = 1 / (1 + np.exp(-(-30.883 + 6.269 * np.log(1.5 * 10) + 5.595 * np.log(1.
    → 5 * 10))))
print("The probability of vasoconstriction with v = r = 1.5: ", prob_v)
```

The probability of vasoconstriction with v = r = 1.5: 0.7764865250020363

1.4 (d) Compare with the result of probit model. Which has smaller deviance?

```
[11]: import theano.tensor as tsr
      from collections import OrderedDict
```

```
[12]: with pm.Model() as probit_model:
      # priors
      intercept = pm.Flat("intercept")
      beta0 = pm.Normal("beta0", mu=0, tau=1.0E-6)
      beta1 = pm.Normal('beta1', mu=0, tau=1.0E-6)

      # linear predictor
      theta_p = intercept + beta0 * df["air_log"] + beta1 * df["rate_log"]

      # Probit transform
      def probit_phi(x):
          mu = 0
          sd = 1
          return 0.5 * (1 + tsr.erf((x - mu) / (sd * tsr.sqrt(2))))

      theta = probit_phi(theta_p)

      # likelihood
      y = pm.Bernoulli('y', p=theta, observed=df["vasoconstriction"])

      trace_prob = pm.sample(5000, tune=5000, init='adapt_diag')
      posterior_predictive_prob = pm.sample_posterior_predictive(trace_prob)
```

Auto-assigning NUTS sampler...

Initializing NUTS using adapt_diag...

ERROR (theano.gof.opt): Optimization failure due to: local_grad_log_erfc_neg

ERROR (theano.gof.opt): node:

Elemwise{true_div,no_inplace}(Elemwise{mul,no_inplace}.0,

Elemwise{erfc,no_inplace}.0)

ERROR (theano.gof.opt): TRACEBACK:

ERROR (theano.gof.opt): Traceback (most recent call last):

File "/opt/anaconda3/lib/python3.7/site-packages/theano/gof/opt.py", line 2034, in process_node

replacements = lopt.transform(node)

File "/opt/anaconda3/lib/python3.7/site-packages/theano/tensor/opt.py", line 6789, in local_grad_log_erfc_neg

if not exp.owner.inputs[0].owner:

AttributeError: 'NoneType' object has no attribute 'owner'

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [beta1, beta0, intercept]

Sampling 4 chains for 5_000 tune and 5_000 draw iterations (20_000 + 20_000 draws total) took 55 seconds.

There were 518 divergences after tuning. Increase `target_accept` or reparameterize.

The acceptance probability does not match the target. It is 0.6682940921512561, but should be close to 0.8. Try to increase the number of tuning steps.

The acceptance probability does not match the target. It is 0.906901957639155, but should be close to 0.8. Try to increase the number of tuning steps.

The number of effective samples is smaller than 10% for some parameters.

```
[13]: model_trace_dict = {"logistic_model": trace_log, "probit_model": trace_prob}
      az.compare(model_trace_dict, ic='WAIC', scale="deviance")
```

```
[13]:
```

	rank	waic	p_waic	d_waic	weight	se	dse \
probit_model	0	36.107	3.18883	0	0.601672	9.70566	0
logistic_model	1	37.1122	3.9323	1.00522	0.398328	8.57837	1.59664

	warning	waic_scale
probit_model	True	deviance
logistic_model	True	deviance

Based on the above table, we could see the probit model has smaller deviance than the logistic model.

2 Magnesium Ammonium Phosphate and Chrysanthemums

Walpole et al. (2007) provide data from a study on the effect of magnesium ammonium phosphate on the height of chrysanthemums, which was conducted at George Mason University in order to determine a possible optimum level of fertilization, based on the enhanced vertical growth response of the chrysanthemums. Forty chrysanthemum seedlings were assigned to 4 groups, each containing 10 plants. Each was planted in a similar pot containing a uniform growth medium. An increasing concentration of MgNH_4PO_4 , measured in grams per bushel, was added to each plant. The 4 groups of plants were grown under uniform conditions in a greenhouse for a period of 4 weeks.

The treatments and the respective changes in heights, measured in centimeters, are given in the following table:

```
[14]: g50 = np.array([13.2, 12.4, 12.8, 17.2, 13.0, 14.0, 14.2, 21.6, 15.0, 20.0])
      g100 = np.array([16.0, 12.6, 14.8, 13.0, 14.0, 23.6, 14.0, 17.0, 22.2, 24.4])
      g200 = np.array([7.8, 14.4, 20.0, 15.8, 17.0, 27.0, 19.6, 18.0, 20.2, 23.2])
      g400 = np.array([21.0, 14.8, 19.1, 15.8, 18.0, 26.0, 21.1, 22.0, 25.0, 18.2])
```

Solve the problem as a Bayesian one-way ANOVA. Use STZ constraints on treatment effects.

2.1 (a) Do different concentrations of MgNH_4PO_4 affect the average attained height of chrysanthemums? Look at the 95% credible sets for the differences between treatment effects.

Explanation on choices of priors:

In pymc3, if the priors are too non-informative, the sampling will fail because there are many zeros when the program takes derivatives. Therefore, we released the constraints and chose weakly informative priors on the mean, α_2 , α_3 and α_4 .

For sigma, the best scenario is using Inversegamma with $\alpha = 0.0001$ and $\beta = 0.0001$. However, the same thing happened. So we released the constraints, too. We changed α and β to 0.01 and 0.01, respectively.

```
[15]: with pm.Model() as ANOVA:
      mu = pm.Normal("mu", mu=0, sigma=1)
      sigma = pm.InverseGamma("sigma", alpha=0.01, beta=0.01)
      alpha2 = pm.Normal("alpha2", mu=0, sigma=1)
      alpha3 = pm.Normal("alpha3", mu=0, sigma=1)
      alpha4 = pm.Normal("alpha4", mu=0, sigma=1)
      # STZ constraints
      alpha1 = pm.Deterministic("alpha1", -(alpha2+alpha3+alpha4))

      # likelihood
      treatment_50 = pm.Normal("g50", mu=mu+alpha1, sigma=sigma, observed=g50)
      treatment_100 = pm.Normal("g100", mu=mu+alpha2, sigma=sigma, observed=g100)
      treatment_200 = pm.Normal("g200", mu=mu+alpha3, sigma=sigma, observed=g200)
      treatment_400 = pm.Normal("g400", mu=mu+alpha4, sigma=sigma, observed=g400)
```

```
[16]: with ANOVA:
      alpha1_diff_alpha2 = pm.Deterministic('a1-a2', alpha1 - alpha2)
      alpha1_diff_alpha3 = pm.Deterministic('a1-a3', alpha1 - alpha3)
      alpha1_diff_alpha4 = pm.Deterministic('a1-a4', alpha1 - alpha4)

      alpha2_diff_alpha3 = pm.Deterministic('a2-a3', alpha2 - alpha3)
      alpha2_diff_alpha4 = pm.Deterministic('a2-a4', alpha2 - alpha4)
      alpha3_diff_alpha4 = pm.Deterministic('a3-a4', alpha3 - alpha4)
```

```
[17]: with ANOVA:
      trace_anova = pm.sample(5000, tune=5000, init='adapt_diag')
```

Auto-assigning NUTS sampler...

Initializing NUTS using adapt_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [alpha4, alpha3, alpha2, sigma, mu]

Sampling 4 chains for 5_000 tune and 5_000 draw iterations (20_000 + 20_000 draws total) took 30 seconds.

```
[18]: az.summary(trace_anova, stat_funcs={"hdi_2.5%": lambda x: np.percentile(x, 2.5),
      ↪ "hdi_97.5%": lambda x : np.percentile(x, 97.5)})
```

```
[18]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_mean	ess_sd \
mu	2.427	1.085	0.378	4.434	0.008	0.006	17454.0	16639.0
alpha2	0.057	0.972	-1.689	1.970	0.006	0.007	27958.0	10118.0
alpha3	0.102	0.962	-1.654	1.953	0.006	0.007	27518.0	10682.0
alpha4	0.158	0.964	-1.661	1.958	0.006	0.007	30313.0	10545.0
sigma	16.241	2.138	12.448	20.348	0.016	0.012	17266.0	17266.0
alpha1	-0.317	1.601	-3.270	2.760	0.009	0.011	30737.0	11046.0
a1-a2	-0.374	2.285	-4.663	3.914	0.013	0.016	30231.0	10665.0
a1-a3	-0.419	2.274	-4.599	3.991	0.013	0.015	29452.0	11092.0
a1-a4	-0.475	2.287	-4.822	3.769	0.013	0.016	31153.0	10730.0
a2-a3	-0.045	1.409	-2.618	2.672	0.009	0.010	26664.0	10495.0
a2-a4	-0.101	1.393	-2.652	2.561	0.009	0.010	26810.0	10476.0
a3-a4	-0.055	1.384	-2.607	2.570	0.008	0.010	28759.0	10374.0

	ess_bulk	ess_tail	r_hat	hdi_2.5%	hdi_97.5%
mu	17461.0	14601.0	1.0	0.321	4.553
alpha2	27981.0	15530.0	1.0	-1.856	1.959
alpha3	27485.0	15232.0	1.0	-1.766	1.996
alpha4	30314.0	15572.0	1.0	-1.760	2.033
sigma	17022.0	14461.0	1.0	12.538	20.935
alpha1	30758.0	15896.0	1.0	-3.440	2.841
a1-a2	30239.0	15848.0	1.0	-4.849	4.110
a1-a3	29465.0	15846.0	1.0	-4.925	4.059
a1-a4	31199.0	15740.0	1.0	-4.911	4.056
a2-a3	26658.0	15470.0	1.0	-2.813	2.714
a2-a4	26801.0	15545.0	1.0	-2.809	2.621
a3-a4	28798.0	15198.0	1.0	-2.734	2.654

Do different concentrations of MgNH_4PO_4 affect the average attained height of chrysanthemums?

Based on the 95% credible set (*hdi*2.5% and *hdi*97.5%) in the above table, we could see the differences between alphas. These differences all cover 0 which means different treatments more likely do not affect the average attained height of chrysanthemums.

2.2 (b) Find the 95% credible set for the contrast $\mu_1 - \mu_2 - \mu_3 + \mu_4$

```
[19]: with ANOVA:
      mu_differences = pm.Deterministic("123+4", (mu+alpha1) - (mu+alpha2) -
      ↪ (mu+alpha3) + (mu+alpha4) )
      trace_mu_differences = pm.sample(5000, tune=5000, init='adapt_diag')
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using adapt_diag...
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [alpha4, alpha3, alpha2, sigma, mu]
```


Sampling 4 chains for 5_000 tune and 5_000 draw iterations (20_000 + 20_000 draws total) took 32 seconds.

```
[20]: az.summary(trace_mu_differences, stat_funcs={"hdi_2.5%": lambda x:np.
      ↪percentile(x, 2.5), "hdi_97.5%": lambda x : np.percentile(x, 97.5)})
```

```
[20]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_mean \
mu	2.431	1.070	0.479	4.483	0.008	0.006	18756.0
alpha2	0.052	0.971	-1.748	1.897	0.006	0.007	27182.0
alpha3	0.104	0.978	-1.703	1.947	0.006	0.007	26466.0
alpha4	0.162	0.971	-1.652	1.978	0.007	0.006	20756.0
sigma	16.228	2.131	12.434	20.276	0.016	0.011	17872.0
alpha1	-0.318	1.634	-3.391	2.697	0.010	0.011	26951.0
a1-a2	-0.370	2.324	-4.827	3.944	0.014	0.016	27645.0
a1-a3	-0.423	2.329	-4.769	3.881	0.014	0.016	27860.0
a1-a4	-0.481	2.315	-4.933	3.740	0.015	0.015	24277.0
a2-a3	-0.052	1.392	-2.641	2.575	0.009	0.010	26296.0
a2-a4	-0.110	1.395	-2.741	2.486	0.009	0.010	22736.0
a3-a4	-0.058	1.404	-2.758	2.511	0.009	0.010	21936.0
mu1-mu2-mu3+mu4	-0.312	2.730	-5.334	4.952	0.017	0.019	27359.0

	ess_sd	ess_bulk	ess_tail	r_hat	hdi_2.5%	hdi_97.5%
mu	17889.0	18746.0	16089.0	1.0	0.331	4.497
alpha2	9812.0	27181.0	15304.0	1.0	-1.862	1.946
alpha3	10612.0	26470.0	14907.0	1.0	-1.800	2.007
alpha4	11464.0	20765.0	15920.0	1.0	-1.721	2.068
sigma	17872.0	17655.0	14169.0	1.0	12.573	20.920
alpha1	10936.0	26919.0	15512.0	1.0	-3.510	2.845
a1-a2	10479.0	27633.0	15421.0	1.0	-4.955	4.153
a1-a3	10827.0	27869.0	15191.0	1.0	-4.938	4.088
a1-a4	11300.0	24279.0	15715.0	1.0	-5.008	4.065
a2-a3	10330.0	26357.0	15717.0	1.0	-2.786	2.668
a2-a4	10619.0	22738.0	15147.0	1.0	-2.866	2.598
a3-a4	10874.0	21939.0	15360.0	1.0	-2.786	2.702
mu1-mu2-mu3+mu4	10512.0	27357.0	14806.0	1.0	-5.659	5.032

```
[43]: print("95% credible set of mu1-mu2-mu3+mu4 : [{}, {}]"
```

```
format(-5.659, 5.032) )
```

95% credible set of mu1-mu2-mu3+mu4 : [-5.659 , 5.032]

This scenario is to test $H_0 : \mu_1 + \mu_4 = \mu_3 + \mu_2 = 0$ VS. $H_1 : \mu_1 + \mu_4 \neq \mu_2 + \mu_3$

$(\mu_1 + \mu_4) - (\mu_2 + \mu_3)$ still covers 0. So two treatments which are added together more likely do not show difference of heights.

3 Hocking–Pendleton Data

This popular data set was constructed by Hocking and Pendelton (1982) to illustrate influential and outlier observations in regression. The data are organized as a matrix of size 26×4 ; the predictors x_1 , x_2 , and x_3 are the first three columns, and the response y is the fourth column. The data are given in `hockpend.dat`

```
[22]: df3 = pd.read_csv("hockpend.dat", sep="\t", header=None).rename(columns={0:"x1", 1:"x2", 2:"x3", 3:"y"})
```

```
[23]: df3.head()
```

```
[23]:
```

	x1	x2	x3	y
0	12.980	0.317	9.998	57.702
1	14.295	2.028	6.776	59.295
2	15.531	5.305	2.947	55.166
3	15.133	4.738	4.201	55.767
4	15.342	7.038	2.053	51.722

3.1 (a) Fit the linear regression model with the three covariates, report the parameter estimates and Bayesian R^2

```
[24]: with pm.Model() as linear_model:
      pm.glm.linear.GLM(y=df3["y"], x= df3[["x1", "x2", "x3"]], intercept=True,
                        family=pm.glm.families.Normal())
      trace_lm = pm.sample(5000, tune=5000, init='adapt_diag')
      posterior_predictive_lm = pm.sample_posterior_predictive(trace_lm)
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using adapt_diag...
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [sd, x3, x2, x1, Intercept]
```

```
Sampling 4 chains for 5_000 tune and 5_000 draw iterations (20_000 + 20_000
draws total) took 90 seconds.
There were 249 divergences after tuning. Increase `target_accept` or
reparameterize.
There were 125 divergences after tuning. Increase `target_accept` or
reparameterize.
There were 1695 divergences after tuning. Increase `target_accept` or
reparameterize.
The acceptance probability does not match the target. It is 0.4309795324404258,
but should be close to 0.8. Try to increase the number of tuning steps.
There were 26 divergences after tuning. Increase `target_accept` or
reparameterize.
The rhat statistic is larger than 1.05 for some parameters. This indicates
```

slight problems during sampling.
The estimated number of effective samples is smaller than 200 for some parameters.

```
[25]: az.summary(trace_lm)
```

```
[25]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_mean	\
Intercept	9.303	9.702	-9.351	25.252	1.588	1.132	37.0	
x1	3.399	0.547	2.505	4.487	0.085	0.060	42.0	
x2	-1.461	0.251	-1.874	-0.972	0.027	0.019	86.0	
x3	0.318	0.280	-0.162	0.840	0.040	0.029	48.0	
sd	2.638	0.438	1.847	3.406	0.043	0.031	102.0	

	ess_sd	ess_bulk	ess_tail	r_hat
Intercept	37.0	37.0	784.0	1.07
x1	42.0	41.0	1411.0	1.06
x2	86.0	86.0	745.0	1.03
x3	48.0	47.0	644.0	1.06
sd	102.0	76.0	47.0	1.03

There are five estimated parameters. The estimated coefficients of the intercept, x_1 , x_2 , x_3 and standard deviation are 9.303, 3.399, -1.461, 0.318 and 2.638, respectively.

```
[26]: r2_scores = []
y_true = df3["y"]
for i in range(len(posterior_predictive_lm["y"])):
    y_pred = posterior_predictive_lm["y"][i]
    r2 = az.r2_score(y_true, y_pred)[0]
    r2_scores.append(r2)
print("Mean of BR2: ", np.mean(r2_scores))
print("Standard deviation of BR2: ", np.std(r2_scores))
```

Mean of BR2: 0.7605379992456223

Standard deviation of BR2: 0.061840177107345136

3.2 (b) Is any of the 26 observations influential or outlier (in the sense of CPO and cumulative)?

3.2.1 Cumulative

We use Cumulative to check whether each data point is an outlier. The concept is to use the distribution defined in each iteration and then we will check where the observed value locates in its cumulative distribution. So after simulation, we could see the means of each data point. If the data point is close to 1 or 0, it is more likely to conclude it is an outlier.

Based on the above samples, our model is:

$$y_i = \beta_0 + \beta_1 * x_{i1} + \beta_2 * x_{i2} + \beta_3 * x_{i3} + \epsilon_i, \text{ where } \epsilon_i \sim N(0, \sigma^2), i = 1, 2, \dots, 26$$

And, we define each coefficient except β_0 as $N(0, 10^{-5})$. As for the intercept, we use the Flat prior.

$$y_i \sim N(\beta_0 + \beta_1 * x_{i1} + \beta_2 * x_{i2} + \beta_3 * x_{i3}, \eta = sd^2)$$

```
[27]: from scipy.stats import norm
```

```
[28]: df_trace = pm.backends.tracetab.trace_to_dataframe(trace_lm)
df_trace.head()
```

```
[28]:
```

	Intercept	x1	x2	x3	sd
0	5.601625	3.622686	-1.233575	0.274603	2.855812
1	7.631317	3.312460	-1.100052	0.625191	2.919152
2	5.730193	3.375129	-0.908699	0.620648	2.695725
3	5.979099	3.553475	-1.339585	0.523389	3.046800
4	3.279882	3.805025	-1.496413	0.331836	2.304360

```
[29]: cuy = np.zeros(26)

for i in range(df_trace.shape[0]):
    intercept = df_trace.iloc[i, 0]
    b1 = df_trace.iloc[i, 1]
    b2 = df_trace.iloc[i, 2]
    b3 = df_trace.iloc[i, 3]
    sd = df_trace.iloc[i, 4]
    for j in range(26):
        obs = df3["y"][j]
        cuy[j] += norm.cdf(obs, loc=intercept + b1 * df3["x1"][j] + b2 *
→df3["x2"][j] + b3 * df3["x3"][j], scale=sd)
```

```
[30]: outlier_check = cuy / df_trace.shape[0]
```

```
[31]: outlier_check
```

```
[31]: array([0.70828015, 0.79224846, 0.48248859, 0.5920942 , 0.4845708 ,
          0.72007163, 0.69536282, 0.39574399, 0.66401189, 0.69294881,
          0.22093436, 0.52572268, 0.62230292, 0.3701755 , 0.00388224,
          0.36372404, 0.94331224, 0.03018092, 0.30486989, 0.77054044,
          0.53216637, 0.68297244, 0.56086072, 0.5843513 , 0.47637694,
          0.40871561])
```

```
[32]: outlier_check[outlier_check<0.1]
```

```
[32]: array([0.00388224, 0.03018092])
```

```
[33]: outlier_check[outlier_check>0.9]
```

```
[33]: array([0.94331224])
```

```
[45]: np.where(outlier_check==0.0038822393888119306)[0]+1
```

```
[45]: array([15])
```

```
[46]: np.where(outlier_check== 0.030180919979533604)[0]+1
```

```
[46]: array([18])
```

```
[47]: np.where(outlier_check==0.9433122366137262)[0]+1
```

```
[47]: array([17])
```

Originally, I chose the $\alpha = 0.2$ and found the 15th, 17th and 18th observations are more likely outliers since the means of these observations are either too close to 0 and 1.

If we choose the smaller confidence level, like we change α to $\alpha = 0.05$, the obvious outlier is the 15th data point.

3.3 (c) Find the mean response and prediction response for a new observation with covariates $x_1^* = 10$, $x_2^* = 5$, and $x_3^* = 5$. Report the corresponding 95% credible sets

```
[37]: new_df = pd.DataFrame({"x1": [10], "x2": [5], "x3": [5], "y": [4.813 + 3.648 * 10_000 + (-1.400) * 5 + 0.420 * 5]})
```

```
[38]: with pm.Model() as lm_pred:
    pm.glm.linear.GLM(y=new_df["y"],
                      x=new_df[["x1", "x2", "x3"]], intercept=True,
                      priors={'Intercept': pm.Normal.dist(mu=0, sigma=10.0)},
                      family=pm.glm.families.Normal())
    trace_lm_p = pm.sample(5000, tune=5000, init='adapt_diag')
    df = pm.trace_to_dataframe(trace_lm, include_transformed=True)
    ppc = pm.sample_posterior_predictive(trace=df,
    →to_dict('records'), samples=len(df))
```

Auto-assigning NUTS sampler...

Initializing NUTS using adapt_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [sd, x3, x2, x1, Intercept]

Sampling 4 chains for 5_000 tune and 5_000 draw iterations (20_000 + 20_000 draws total) took 39 seconds.

There were 2345 divergences after tuning. Increase `target_accept` or reparameterize.

The acceptance probability does not match the target. It is 0.3403953364767104, but should be close to 0.8. Try to increase the number of tuning steps.

There were 1798 divergences after tuning. Increase `target_accept` or

reparameterize.

The acceptance probability does not match the target. It is 0.44437999849911075, but should be close to 0.8. Try to increase the number of tuning steps.

There were 3543 divergences after tuning. Increase `target_accept` or reparameterize.

The acceptance probability does not match the target. It is 0.05906703653418209, but should be close to 0.8. Try to increase the number of tuning steps.

There were 2760 divergences after tuning. Increase `target_accept` or reparameterize.

The acceptance probability does not match the target. It is 0.30560815998471047, but should be close to 0.8. Try to increase the number of tuning steps.

The rhat statistic is larger than 1.4 for some parameters. The sampler did not converge.

The estimated number of effective samples is smaller than 200 for some parameters.

Prediction response:

This means the result is derived based on inputs and coefficients.

```
[1]: 4.813 + 3.648 * 10 + (-1.400) * 5 + 0.420 * 5
```

```
[1]: 36.393000000000001
```

Mean response:

This means that based on each iteration, we get estimated coefficients and then have multiple response values. We take expectation on these values.

```
[40]: ppc["y"].mean()
```

```
[40]: 37.576426132208155
```

```
[41]: az.hdi(ppc["y"], hdi_prob=0.95)
```

```
/opt/anaconda3/lib/python3.7/site-packages/arviz/stats/stats.py:487:
```

```
FutureWarning: hdi currently interprets 2d data as (draw, shape) but this will  
change in a future release to (chain, draw) for coherence with other functions  
FutureWarning,
```

```
[41]: array([[29.97973519, 44.50649613]])
```

The prediction response is 36.393000000000001.

The mean response is 37.576426132208155 and the 95% credible set is between 29.97973519 and 44.50649613.